

Exercises week 2

Last update: 2022/09/08

Exercise 2.1 Consider the Readers and Writers problem we saw in class. As a reminder, here is the specification of the problem:

- Several reader and writer threads want to access a shared resource.
- Many readers may access the resource at the same time as long as there are no writers.
- At most one writer may access the resource if and only if there are no readers.

Mandatory

1. Use Java Intrinsic Locks (i.e., `synchronized`) to implement a monitor ensuring that the access to the shared resource by reader and writer threads is according to the specification above. You may use the code in the lectures and Chapter 8 of Herlihy as inspiration, but do not feel obliged copy that structure.
2. Is your solution fair towards writer threads? In other words, does your solution ensure that if a writer thread wants to write, then it will eventually do so? If so, explain why. If not, modify part 1. so that your implementation satisfies this fairness requirement, and then explain why your new solution satisfies the requirement.

Challenging

3. What type of fairness does your solution enforce? Hint: Consider the two types of fairness we discussed in class, see lecture slides.
4. Is it possible to ensure strong fairness using `ReentrantLock` or intrinsic java locks (`synchronized`)? Explain why.

Exercise 2.2 Consider the lecture's example in file `TestMutableInteger.java`, which contains this definition of class `MutableInteger`:

```
// WARNING: Not ready for usage by concurrent programs
class MutableInteger {
    private int value = 0;
    public void set(int value) {
        this.value = value;
    }
    public int get() {
        return value;
    }
}
```

For instance, as mentioned in Goetz, this class cannot be used to reliably communicate an integer from one thread to another, as attempted here:

```
final MutableInteger mi = new MutableInteger();
Thread t = new Thread(() -> {
    while (mi.get() == 0)          // Loop while zero
    { /* Do nothing*/ }
    System.out.println("I completed, mi = " + mi.get());
});
t.start();
try { Thread.sleep(500); } catch (InterruptedException e) {
    e.printStackTrace(); }
```

```
mi.set(42);
System.out.println("mi set to 42, waiting for thread ...");
try { t.join(); } catch (InterruptedException e) { e.printStackTrace(); }
System.out.println("Thread t completed, and so does main");
```

Mandatory

1. Compile and run the example as is. Do you observe the "main" thread's write to `mi.value` remains invisible to the `t` thread, so that it loops forever? Independently of your observation, is it possible that the program loops forever? Explain your answer.
2. Use Java Intrinsic Locks (`synchronized`) on the methods of the `MutableInteger` to ensure that thread `t` always terminates. Explain why your solution prevents thread `t` from running forever.
3. Would thread `t` always terminate if `get()` is not defined as `synchronized`? Explain your answer.
4. Remove all the locks in the program, and define `value` in `MutableInteger` as a `volatile` variable. Does thread `t` always terminate in this case? Explain your answer.

Challenging

5. Explain parts 3. and 4. in terms of the *happens-before* relation.

Exercise 2.3 Consider the small artificial program in file `TestLocking0.java`. In class `Mystery`, the single mutable field `sum` is `private`, and all methods are `synchronized`, so superficially the class seems that no concurrent sequence of method calls can lead to race conditions.

Mandatory

1. Compile the program and run it several times. Show the results you get. Are there any race conditions?
2. Explain why race conditions appear when `t1` and `t2` use the `Mystery` object. Hint: Consider (a) what it means for an instance method to be `synchronized`, and (b) what it means for a static method to be `synchronized`.
3. Implement a new version of the class `Mystery` so that the execution of `t1` and `t2` does not produce race conditions, *without* changing the modifiers of the field and methods in the `Mystery` class. That is, you should not make any static field into an instance field (or vice versa), and you should not make any static method into an instance method (or vice versa).

Explain why your new implementation does not have race conditions.

Exercises week 1

Last update: 2022/09/01

Practical info before you start

Abbreviations

The following abbreviations are used in the exercise sheets:

- “Goetz” means Goetz *et al.*, *Java Concurrency in Practice*, Addison-Wesley 2006.
- “Herlihy” means Herlihy *et al.*, *The Art of Multiprocessor Programming*, Morgan Kaufmann, 2020.

The exercises are a way for you to get a practical understanding of the material. In addition, they serve as the outset for bi-weekly feedback.

Oral Feedback

Feedback is given orally. Please see the document in our Github repository (<https://github.itu.dk/jst/PCPP2022-public/blob/main/general-info/assignment-submissions-and-oral-feedback.md>) regarding how to book an oral feedback session, and how feedback sessions are conducted.

Mandatory/Challenging Exercises

Exercises are labelled as *mandatory* or *challenging*. For a submission to be considered accepted, all mandatory exercises must be successfully completed (and approved). Furthermore, the code of mandatory exercises, shown during oral sessions, must compile and run. Code that is close to correct, but that does not compile and run will not be accepted; note again that this is only for mandatory assignments.

We acknowledge that different students strive to strike different balances between the different classes and between study and other aspects of their lives. Consequently, we include *challenging* exercises. These exercises are *optional*, they are not required for acceptance. However, completing the challenging exercises increases your chances of obtaining high marks at the exam.

Groups should be composed by students at the same ambition level.

Installing JDK/Gradle and Running Exercise Code

In this course, we will use the Gradle tool and Java JDK 8 or higher for the exercises (and lectures). We have provided a guide on how to set up the programming environment we will use in the course. The guide is in our Github repository (<https://github.itu.dk/jst/PCPP2022-public/blob/main/general-info/guide-using-gradle-for-exercises.md>). Please follow the guide to ensure that you set up the programming environment in your machine. The guide provides information for different operating systems. If you have any trouble on getting things running please contact us, e.g., by attending exercise sessions or posting in the Questions and Answer Forum in LearnIT.

Exercise 1.1 Consider the code in the file `TestLongCounterExperiments.java`. Note that this file contains a class, `LongCounter`, used by two threads:

```
class LongCounter {
    private long count = 0;
    public void increment() {
        count = count + 1;
    }
    public long get() {
        return count;
    }
}
```

Mandatory

1. The `main` method creates a `LongCounter` object. Then it creates and starts two threads that run concurrently, and each increments the `count` field 10 million times by calling method `increment`.

What output values do you get? Do you get the expected output, i.e., 20 million?

2. Reduce the `counts` value from 10 million to 100, recompile, and rerun the code. It is now likely that you get the expected result (200) in every run.

Explain how this could be. Is it guaranteed that the output is always 200?

3. The `increment` method in `LongCounter` uses the assignment

```
count = count + 1;
```

to add one to `count`. This could be expressed also as `count += 1` or as `count++`.

Do you think it would make any difference to use one of these forms instead? Why? Change the code and run it. Do you see any difference in the results for any of these alternatives?

4. Set the value of `counts` back to 10 million. Use Java `ReentrantLock` to ensure that the output of the program equals 20 million. Explain why your solution is correct, and why no other output is possible.

Note: In your explanation, please use the concepts and vocabulary introduced during the lecture, e.g., critical sections, interleavings, race conditions, mutual exclusion, etc.

Note II: The notes above applies to all exercises asking you to explain the correctness of your solution.

5. By using the `ReentrantLock` in the exercise above, you have defined a *critical section*. Does your critical section contain the least number of lines of code? If so, explain why. If not, fix it and explain why your new critical section contains the least number of lines of code.

Hint: Recall that the critical section should only include the parts of the program that only one thread can execute at the same time.

Challenging

6. Decompile the methods `increment` from above to see the byte code in the three versions (as is, `+=`, `++`). The basic decompiler is `javap`. The decompiler takes as input a (target) `.class` file. In Gradle projects, `.class` files are located in the directory `app/build/classes/`—after compiling the `.java` files. The flag `-c` decompiles the code of a class. Does the output of `javap` verify or refuse the explanation you provided in part 3.?
7. Extend the `LongCounter` class with a `decrement()` method which subtracts 1 from the `count` field without using locks. Change the code in `main` so that `t1` calls `decrement` 10 million times, and `t2` calls `increment` 10 million times, on a `LongCounter` instance. What should the expected output be after both threads have completed?

Use `ReentrantLock` to ensure that the program outputs 0. Explain why your solution is correct, and why no other output is possible.

8. Explain, in terms of the *happens-before* relation, why your solutions for part 4. and 7. produce the expected output.
9. Remove the `ReentrantLock` you added for parts 4. and 7. Set the variable `counts` to 3.

What is the minimum value that the program prints? Does the minimum value change if we set `counts` to a larger number (e.g., 20 million)?

Provide an interleaving showing how the minimum value output can occur.

Exercise 1.2 Consider this class, whose `print` method prints a dash “-”, waits for 50 milliseconds, and then prints a vertical bar “|”:

```
class Printer {
    public void print() {
        System.out.print("-");
        try { Thread.sleep(50); } catch (InterruptedException exn) { }
        System.out.print("|");
    }
}
```

Mandatory

1. Write a program that creates a `Printer` object `p`, and then creates and starts two threads. Each thread must call `p.print()` forever. Note: You can easily run your program using the gradle project for the exercises by placing your code in the directory `week01exercises/app/src/main/java/exercises01/` (remember to add `package exercises01;` as the first line of your Java files).

You will observe that, most of the time, your program print the dash and bar symbols alternate neatly as in `-|-|-|-|-|-|-|`. But occasionally two bars are printed in a row, or two dashes are printed in a row, creating small “weaving faults” like those shown below:

[illegible]

2. Describe and provide an interleaving where this happens.
3. Use Java `ReentrantLock` to ensure that the program outputs the expected sequence `-|-|-|...`

Compile and run the improved program to see whether it works. Explain why your solution is correct, and why it is not possible for incorrect patterns, such as in the output above, to appear.

Challenging

4. Explain, in terms of the *happens-before* relation, why your solution for part 3 produces the correct output.

Exercise 1.3 Imagine that due to the COVID-19 pandemic Tivoli decides to limit the number of visitors to 15000. To this end, you are asked to modify the code for the turnstiles we saw in the lecture. The file `CounterThreads2Covid.java` includes a new constant `MAX_PEOPLE_COVID` equal to 15000. However, the two threads simulate 20000 people entering to the park, so unfortunately some people will not get in : ' (.

Mandatory

1. Modify the behaviour of the `Turnstile` thread class so that that exactly 15000 enter the park; no less no more. To this end, simply add a check before incrementing `counter` to make sure that there are less than 15000 people in the park. Note that the class does not use any type of locks. You must use `ReentrantLock` to ensure that the program outputs the correct value, 15000.
2. Explain why your solution is correct, and why it always output 15000.

Exercise 1.4 In *Goetz* chapter 1.1, three motivations for concurrency is given: resource utilization, fairness and convenience. In the note about concurrency there is an alternative characterization of the different motivations for concurrency (coined by the Norwegian computer scientist Kristen Nygaard):

- **Inherent** User interfaces and other kinds of input/output.
- **Exploitation** Hardware capable of simultaneously executing multiple streams of statements. A special (but important) case is communication and coordination of independent computers on the internet,
- **Hidden** Enabling several programs to share some resources in a manner where each can act as if they had sole ownership.

Mandatory

Neither of the definitions is very precise, so for this exercise, there are many possible and acceptable answers.

1. Compare the categories in the concurrency note (<https://github.itu.dk/jst/PCPP2022-public/blob/main/week01/concurrencyNotes/concurrencyPCPP.pdf>) and *Goetz*, try to find some examples of systems which are included in the categories of *Goetz*, but not in those in the concurrency note, and vice versa (if possible - if not possible, argue why).
2. Find examples of 3 systems in each of the categories in the Concurrency note which you have used yourself (as a programmer or user).